

04_exercise_solution

November 9, 2021

1 Exercise: Diabetes Model

In this exercise, we're going to take the knowledge we gained from the lesson and apply it to the [Diabetes dataset](#). This well known dataset already has it's data cleaned and normalized, so no need to do any of those steps. The steps required to complete this exercise are as follows:

1. Load the diabetes dataset into a dataframe.
2. Check the table summary to show that indeed the mean is zero for all features.
3. Split the dataset into train, validation, and test sets
4. Use a linear regression Ridge model to fit and score:
 1. Fit and score on the whole dataset
 2. Fit on train, score on validation, using default model
 3. Fit on train, score on validation, using hyperparameters model
 4. Fit on train, score on test, using hyperparameterized model
5. Plot all scores in a bar graph

```
In [1]: import numpy as np
import pandas as pd
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge
```

```
In [2]: # Load in the diabetes dataset
diabetes = datasets.load_diabetes()
```

```
In [3]: # Create the diabetes `data` dataset as a dataframe and name the columns with `feature_`
df = pd.DataFrame(diabetes["data"], columns=diabetes["feature_names"])

# Include the target as well
df["target"] = diabetes["target"]
```

```
In [4]: df.head()
```

```
Out[4]:
```

	age	sex	bmi	bp	s1	s2	s3	\
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401	
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412	
2	0.085299	0.050680	0.044451	-0.005671	-0.045599	-0.034194	-0.032356	

```

3 -0.089063 -0.044642 -0.011595 -0.036656 0.012191 0.024991 -0.036038
4 0.005383 -0.044642 -0.036385 0.021872 0.003935 0.015596 0.008142

```

```

      s4      s5      s6 target
0 -0.002592 0.019908 -0.017646 151.0
1 -0.039493 -0.068330 -0.092204 75.0
2 -0.002592 0.002864 -0.025930 141.0
3 0.034309 0.022692 -0.009362 206.0
4 -0.002592 -0.031991 -0.046641 135.0

```

```

In [5]: # Describe df using table summary.
        # No need to normalize, near zero mean.
        df.describe()

```

```

Out [5]:
      count      age      sex      bmi      bp      s1 \
count  4.420000e+02  4.420000e+02  4.420000e+02  4.420000e+02  4.420000e+02
mean   -3.639623e-16  1.309912e-16 -8.013951e-16  1.289818e-16 -9.042540e-17
std     4.761905e-02  4.761905e-02  4.761905e-02  4.761905e-02  4.761905e-02
min    -1.072256e-01 -4.464164e-02 -9.027530e-02 -1.123996e-01 -1.267807e-01
25%    -3.729927e-02 -4.464164e-02 -3.422907e-02 -3.665645e-02 -3.424784e-02
50%     5.383060e-03 -4.464164e-02 -7.283766e-03 -5.670611e-03 -4.320866e-03
75%     3.807591e-02  5.068012e-02  3.124802e-02  3.564384e-02  2.835801e-02
max     1.107267e-01  5.068012e-02  1.705552e-01  1.320442e-01  1.539137e-01

      count      s2      s3      s4      s5      s6 \
count  4.420000e+02  4.420000e+02  4.420000e+02  4.420000e+02  4.420000e+02
mean    1.301121e-16 -4.563971e-16  3.863174e-16 -3.848103e-16 -3.398488e-16
std     4.761905e-02  4.761905e-02  4.761905e-02  4.761905e-02  4.761905e-02
min    -1.156131e-01 -1.023071e-01 -7.639450e-02 -1.260974e-01 -1.377672e-01
25%    -3.035840e-02 -3.511716e-02 -3.949338e-02 -3.324879e-02 -3.317903e-02
50%    -3.819065e-03 -6.584468e-03 -2.592262e-03 -1.947634e-03 -1.077698e-03
75%     2.984439e-02  2.931150e-02  3.430886e-02  3.243323e-02  2.791705e-02
max     1.987880e-01  1.811791e-01  1.852344e-01  1.335990e-01  1.356118e-01

      count      target
count  442.000000
mean   152.133484
std     77.093005
min     25.000000
25%     87.000000
50%    140.500000
75%    211.500000
max    346.000000

```

```

In [6]: # train: 0.8 / test: 0.2
        df_train, df_test = train_test_split(df, test_size=0.2, random_state=0)

        # train: 0.6 / validation: 0.2

```

```

df_train, df_val = train_test_split(df_train, test_size=0.25, random_state=0)

# Final dataset sizes: train: 0.6, validation: 0.2, test: 0.2

In [7]: # How does the model perform on the entire dataset and default model parameters
reg = Ridge().fit(df[diabetes["feature_names"]], df["target"])
all_df_score = reg.score(df[diabetes["feature_names"]], df["target"])
all_df_score

Out[7]: 0.4512313946799055

In [8]: # How does the model perform on the training dataset and default model parameters
# Remember we use the validation dataset score the model
reg = Ridge().fit(df_train[diabetes["feature_names"]], df_train["target"])
val_df_score = reg.score(df_val[diabetes["feature_names"]], df_val["target"])
val_df_score

Out[8]: 0.3934444316580993

In [9]: # How does the model perform on the training dataset and different model parameters
# Change alpha, solver, and max_iter
reg_h = Ridge(alpha=0.01, solver="saga", max_iter=10000).fit(
    df_train[diabetes["feature_names"]], df_train["target"]
)
val_df_h_score = reg_h.score(df_val[diabetes["feature_names"]], df_val["target"])
val_df_h_score

Out[9]: 0.49381740194998924


In [10]: # Use optimized data on the held out test dataset.
test_df_h_score = reg_h.score(df_test[diabetes["feature_names"]], df_test["target"])
test_df_h_score

Out[10]: 0.3399053725161949

In [39]: # Bar plot of all scores from each model fit: all_df_score, val_df_score, val_df_h_score
pd.Series(
    {
        "all_df_score": all_df_score,
        "val_df_score": val_df_score,
        "val_df_h_score": val_df_h_score,
        "test_df_h_score": test_df_h_score,
    }
).plot(kind="bar", legend=False, title="R2 Score of Ridge Model")

Out[39]: <AxesSubplot:title={'center': 'R2 Score of Ridge Model'}>

```



In []: